

A Log Buffer-Based Flash Translation Layer Using Fully-Associative Sector Translation

SANG-WON LEE

Sungkyunkwan University

DONG-JOO PARK

Soongsil University

TAE-SUN CHUNG

Ajou University

DONG-HO LEE

Hanyang University

SANGWON PARK

Hankook University of Foreign Studies

and

HA-JOO SONG

Pukyong National University

Flash memory is being rapidly deployed as data storage for mobile devices such as PDAs, MP3 players, mobile phones, and digital cameras, mainly because of its low electronic power, nonvolatile storage, high performance, physical stability, and portability. One disadvantage of flash memory is that prewritten data cannot be dynamically overwritten. Before overwriting prewritten data, a time-consuming erase operation on the used blocks must precede, which significantly degrades

This work was supported in part by MIC & IITA through IT Leading R&D Support Project, in part by MIC & IITA through Oversea Post-Doctoral Support Program 2005, in part by the Ministry of Information and Communication, Korea under the ITRC support program supervised by the Institute of Information Technology Assessment, IITA-2005-(C1090-0501-0019), and also supported in part by Seoul R&D Program(10660).

Authors' addresses: Sang-Won Lee, School of Information and Communications Engineering, Sungkyunkwan University, Suwon 440-746, Korea; email: swlee@acm.org; Dong-Joo Park, School of Computing, Soongsil University, Seoul 156-743, Korea; email: djpark@ssu.ac.kr; Tae-Sun Chung, College of Information Technology, Ajou University, Suwon 443-749, Korea; email: tschung@ajou.ac.kr; Dong-Ho Lee, Department of Computer Science and Engineering, Hanyang University, Ansan 426-791, Korea; email: dhlee72@cse.hanyang.ac.kr; Sangwon Park, Information Communication Engineering, Hankook University of Foreign Studies, Yongin 449-791, Korea; email: swpark@hufs.ac.kr; Ha-Joo Song, Division of Electronic, Computer, and Telecommunication, Pukyong National University, Busan 608-737, Korea; email: hagusong@pknu.ac.kr.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or direct commercial advantage and that copies show this notice on the first page or initial screen of a display along with the full citation. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, to republish, to post on servers, to redistribute to lists, or to use any component of this work in other works requires prior specific permission and/or a fee. Permissions may be requested from Publications Dept., ACM, Inc., 2 Penn Plaza, Suite 701, New York, NY 10121-0701 USA, fax +1 (212) 869-0481, or permissions@acm.org. © 2007 ACM 1539-9087/2007/07-ART18 \$5.00 DOI 10.1145/1275986.1275990 <http://doi.acm.org/10.1145/1275986.1275990>

the overall write performance of flash memory. In order to solve this “erase-before-write” problem, the flash memory controller can be integrated with a software module, called “flash translation layer (FTL).” Among many FTL schemes available, the log block buffer scheme is considered to be optimum. With this scheme, a small number of log blocks, a kind of write buffer, can improve the performance of write operations by reducing the number of erase operations. However, this scheme can suffer from low space utilization of log blocks. In this paper, we show that there is much room for performance improvement in the log buffer block scheme, and propose an enhanced log block buffer scheme, called FAST (full associative sector translation). Our FAST scheme improves the space utilization of log blocks using fully-associative sector translations for the log block sectors. We also show empirically that our FAST scheme outperforms the pure log block buffer scheme.

ACM Reference Format:

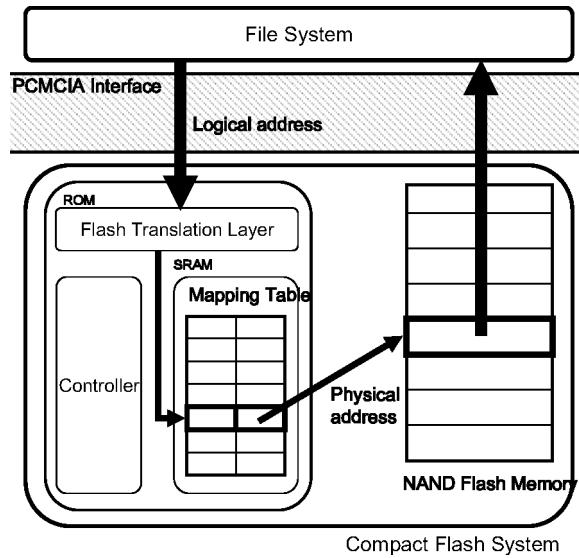
Lee, S.-W., Park, D.-J., Chung, T.-S., Lee, D.-H., Park, S., and Song, H.-J. 2007. A log buffer-based flash translation layer using fully-associative sector translation. *ACM Trans. Embedd. Comput. Syst.* 6, 3, Article 18 (July 2007), 27 pages. DOI = 10.1145/1275986.1275990 <http://doi.acm.org/10.1145/1275986.1275990>

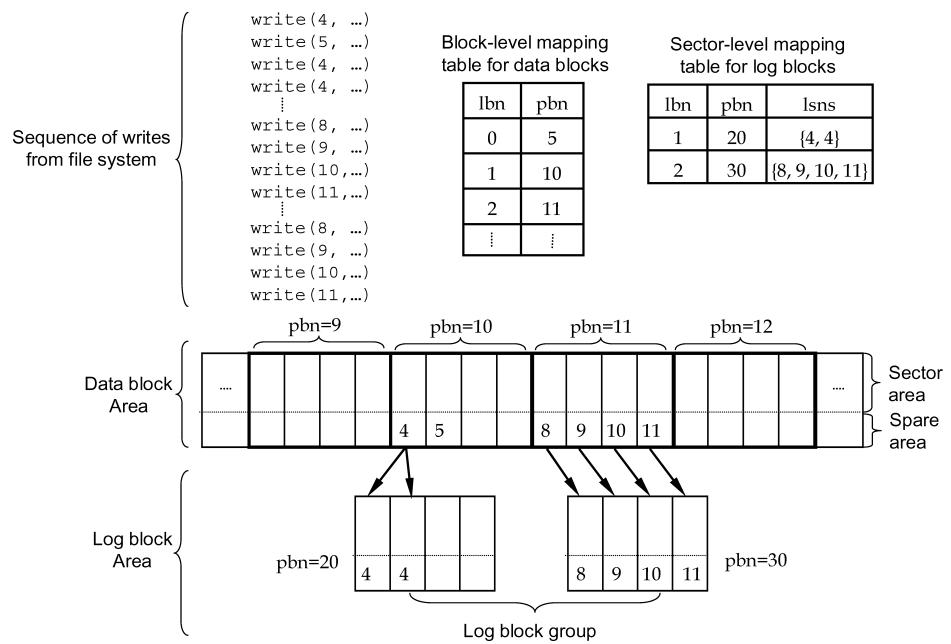
1. INTRODUCTION

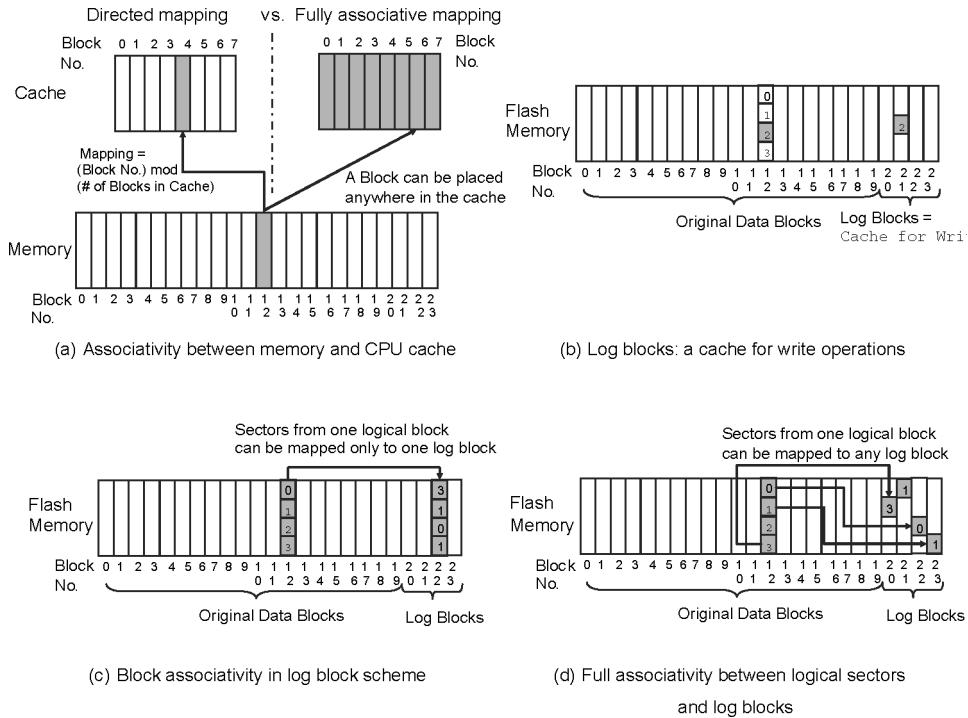
Flash memory is being rapidly deployed as data storage for mobile devices, such as PDAs, MP3 players, mobile phones, and digital cameras, mainly because of its small size, low power consumption, shock resistance, and nonvolatile memory [Douglis et al. 1994; Kim et al. 2002; Gal and Toledo 2005]. Compared to a hard disk with the inevitable mechanical delay in accessing data (that is, seek time and rotational latency), flash memory provides fast uniform random access. For example, the read and write time per sector (typically, 512 bytes) for NAND-flash memory is 15 and 200 μ s [Samsung Electronics 2005], respectively, while each operation in contemporary hard disks takes around 10 ms [Hennessy and Patterson 2003].

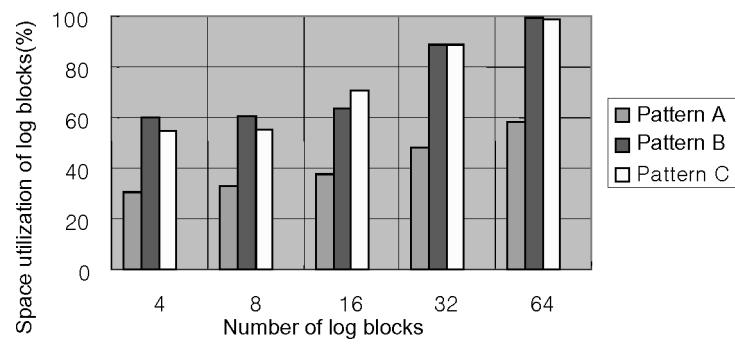
However, flash memory suffers from the write bandwidth problem. A write operation is slower than a read operation by an order of magnitude. In addition, a write operation may have to be preceded by a costly erase operation because flash memory does not allow overwrites. Unfortunately, write operations are performed in a sector unit, while erase operations are executed in a block unit: usually, one block consists of 32 sectors. An erase operation is very time-consuming, compared to a write operation; usually, a per-block erase time is 2 ms [Samsung Electronics 2005]. These inherent characteristics of flash memory reduce write bandwidth, which is the performance bottleneck in flash-based mobile devices.

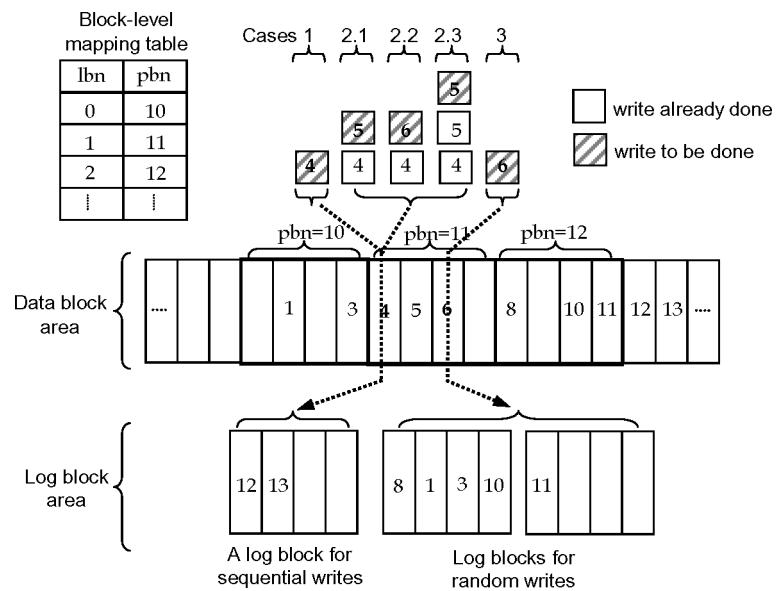
To relieve this performance bottleneck, it is very important to reduce the number of erase operations resulting from write operations. For this, flash memory vendors have adopted an intermediate software module, called flash translation layer (FTL), between the host applications (in general, file systems) and flash memory [Estakhri and Iman 1999; Kim and Lee 2002; Kim et al. 2002;











Algorithm 1 `write(lsn, data)` /* data are logically written to the sector of lsn */

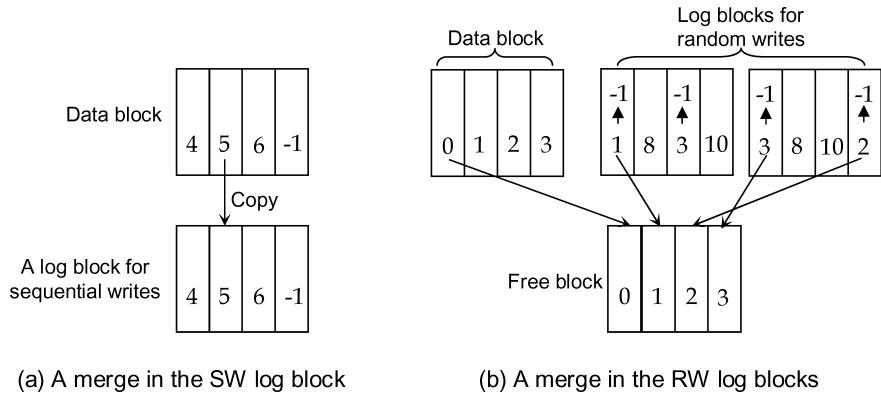
```

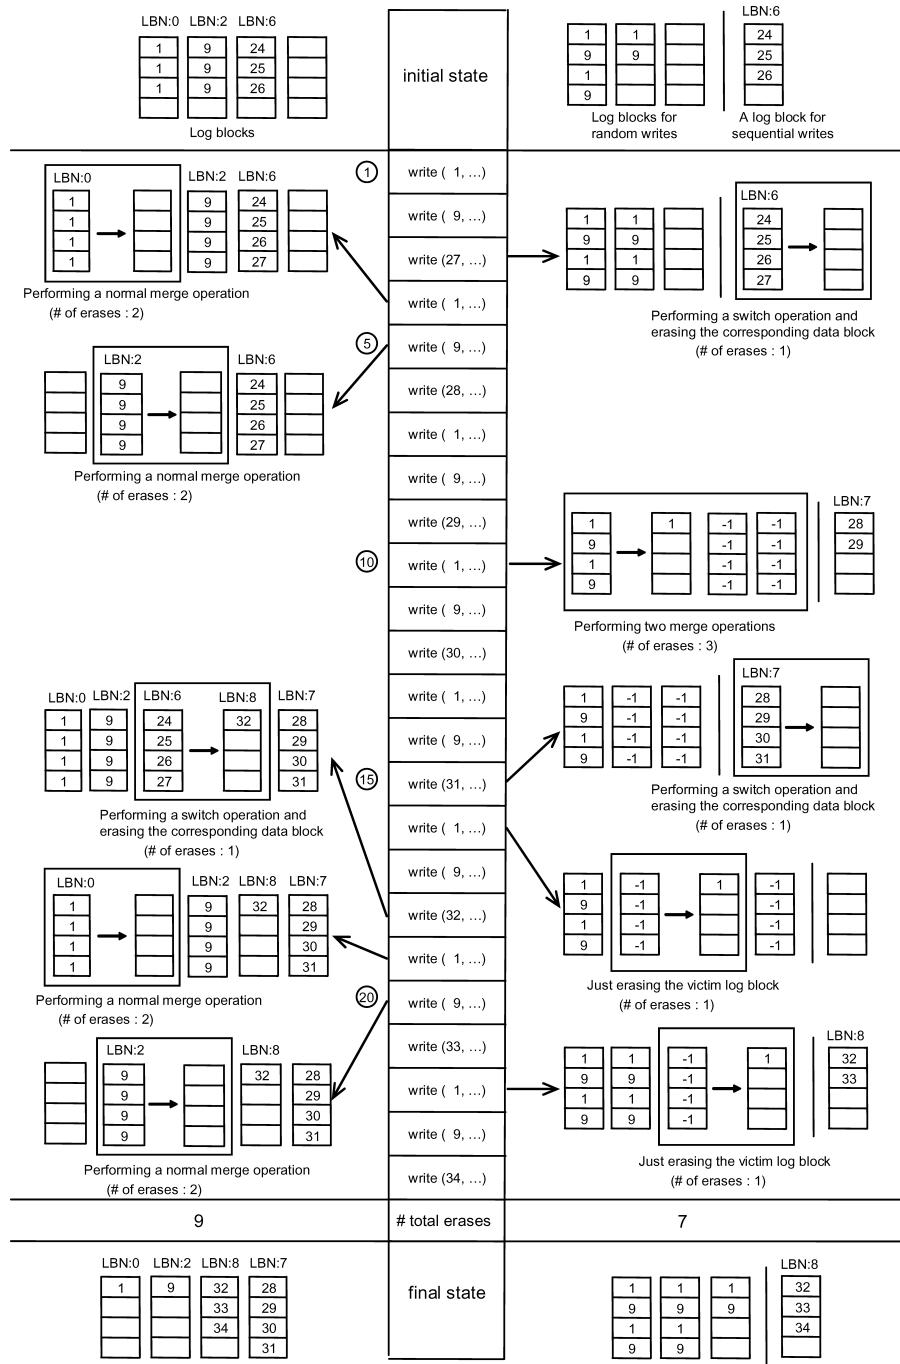
1  lbn := lsn div SectorsPerBlock;
2  offset := lsn mod SectorsPerBlock;
3  pbn := getPbnFromBMT(lbn); /* get pbn from block-level mapping table */
4  if a collision occurs at offset of the data block of pbn
5      call writeToLogblock(lsn, lbn, offset, data);
6  else
7      write data at offset in the data block of pbn;
8  end if
```

Algorithm 2 `writeToLogblock(lsn, lbn, offset, data)`

```

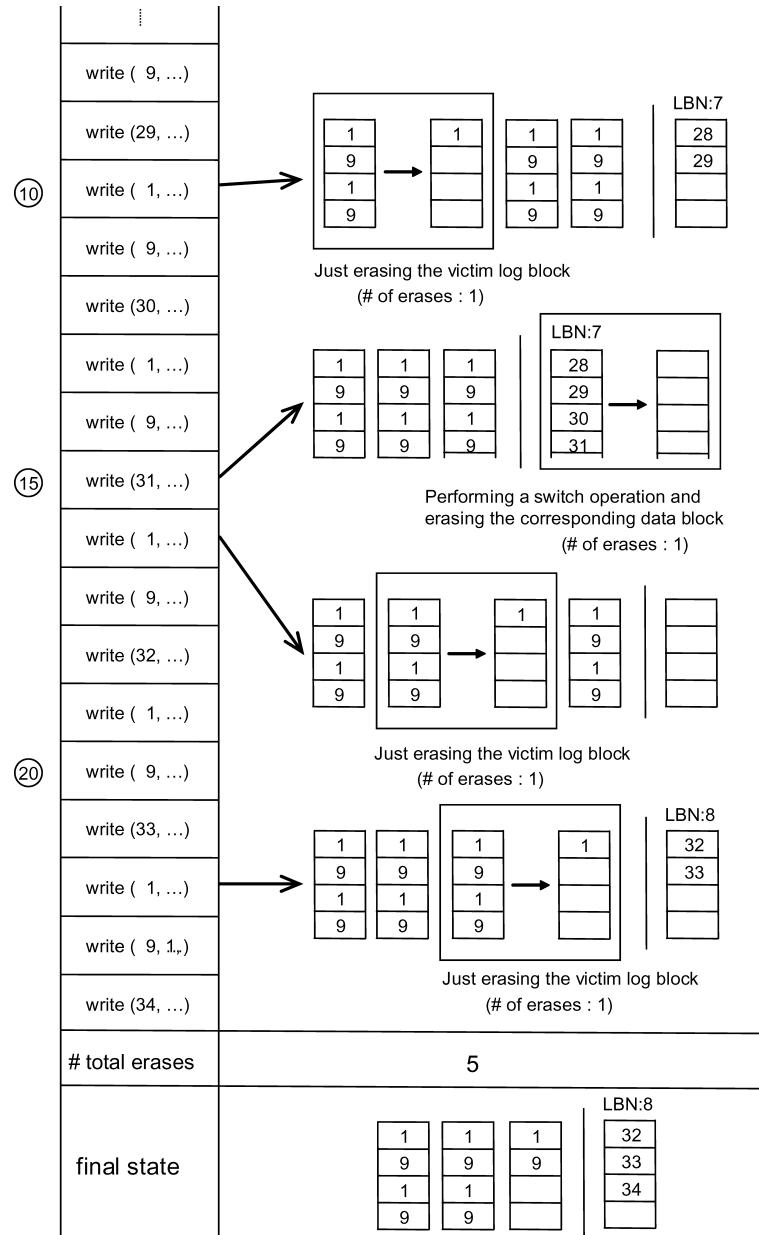
1  if offset is zero /* Case 1 in Figure 5 */
2      if there are no empty sectors in the SW log block
3          /* the log block is filled with sequentially written sectors */
4          perform a switch operation between the SW log block and
5          its corresponding data block;
6          /* after switch, the data block is erased and returned to the free-block list */
7      else
8          /* before merge, a new block is allocated from the free-block list */
9          merge the SW log block with its corresponding data block;
10         /* after merge, the two blocks are erased and returned to the free-block list */
11     end if
12     get a block from the free-block list and use it as a SW log block;
13     append data to the SW log block;
14     update the SW log block part of the sector-mapping table;
15   else
16     if the current owner of the SW log block is the same with lbn
17         last_lsn := getLastLsnFromSMT(lbn); /* SMT: sector-mapping table */
18         if lsn is equivalent with (last_lsn+1) /* Case 2.1 in Figure 5 */
19             append data to the SW log block;
20         else /* Case 2.2 and 2.3 in Figure 5*/
21             merge the SW log block with its corresponding data block;
22             get a block from the free-block list and use it as a SW log block;
23         end if
24         update the SW log block part of the sector-mapping table;
25     else /* Case 3 in Figure 5 */
26         if there are no rooms in the RW log blocks to write data
27             select the first block of the RW log block list as a victim;
28             merge the victim with its corresponding data block;
29             get a block from the free-block list and add it to the end of
30             the RW log block list;
31             update the RW log block part of the sector-mapping table;
32         end if
33         append data to the RW log blocks;
34     end if
35   end if
```





(a) BAST

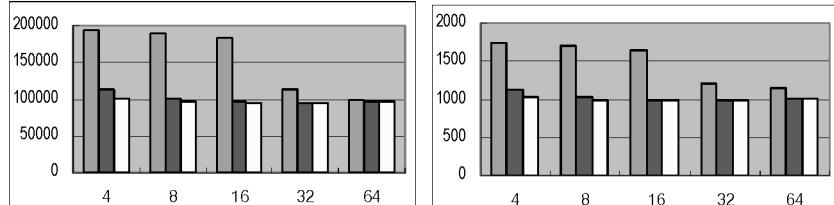
(b) FAST



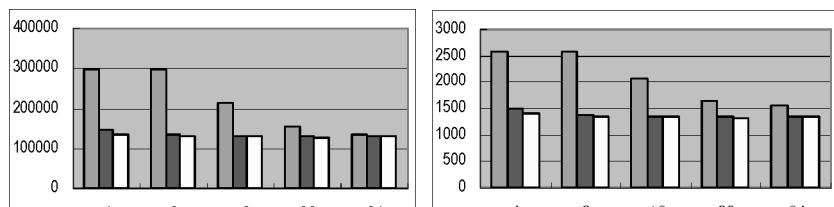


■ BAST ■ FAST □ O-FAST

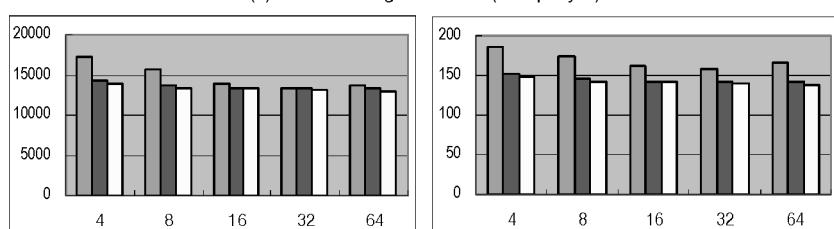
X-axis : # of log blocks, Y-axis in left side : erase count, Y-axis in right side : elapsed time(secs).



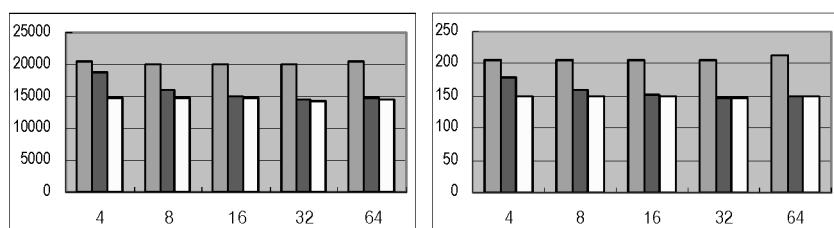
(a) Pattern A: Digital Camera(Company A)



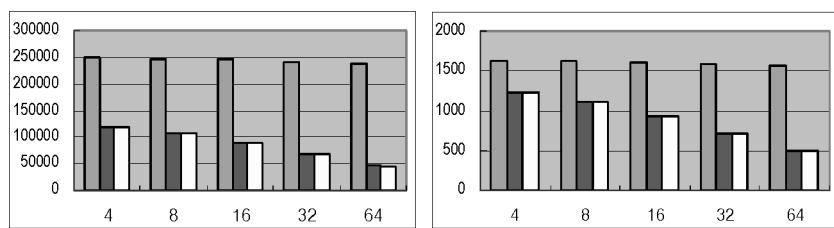
(b) Pattern B: Digital Camera(Company B)



(c) Pattern C: Linux



(d) Pattern D: Symbian



(e) Pattern E: Random

<hr/>			

